

Foundations of High Performance Computing

Lecture 3: HPC software stack

“Foundation of HPC” course



**DATA SCIENCE &
SCIENTIFIC COMPUTING**

2021-2022 Stefano Cozzini

Agenda

A first look of the software stack

Local resource manager: queue system

Scientific software

Compilers

Libraries

From reference 1...

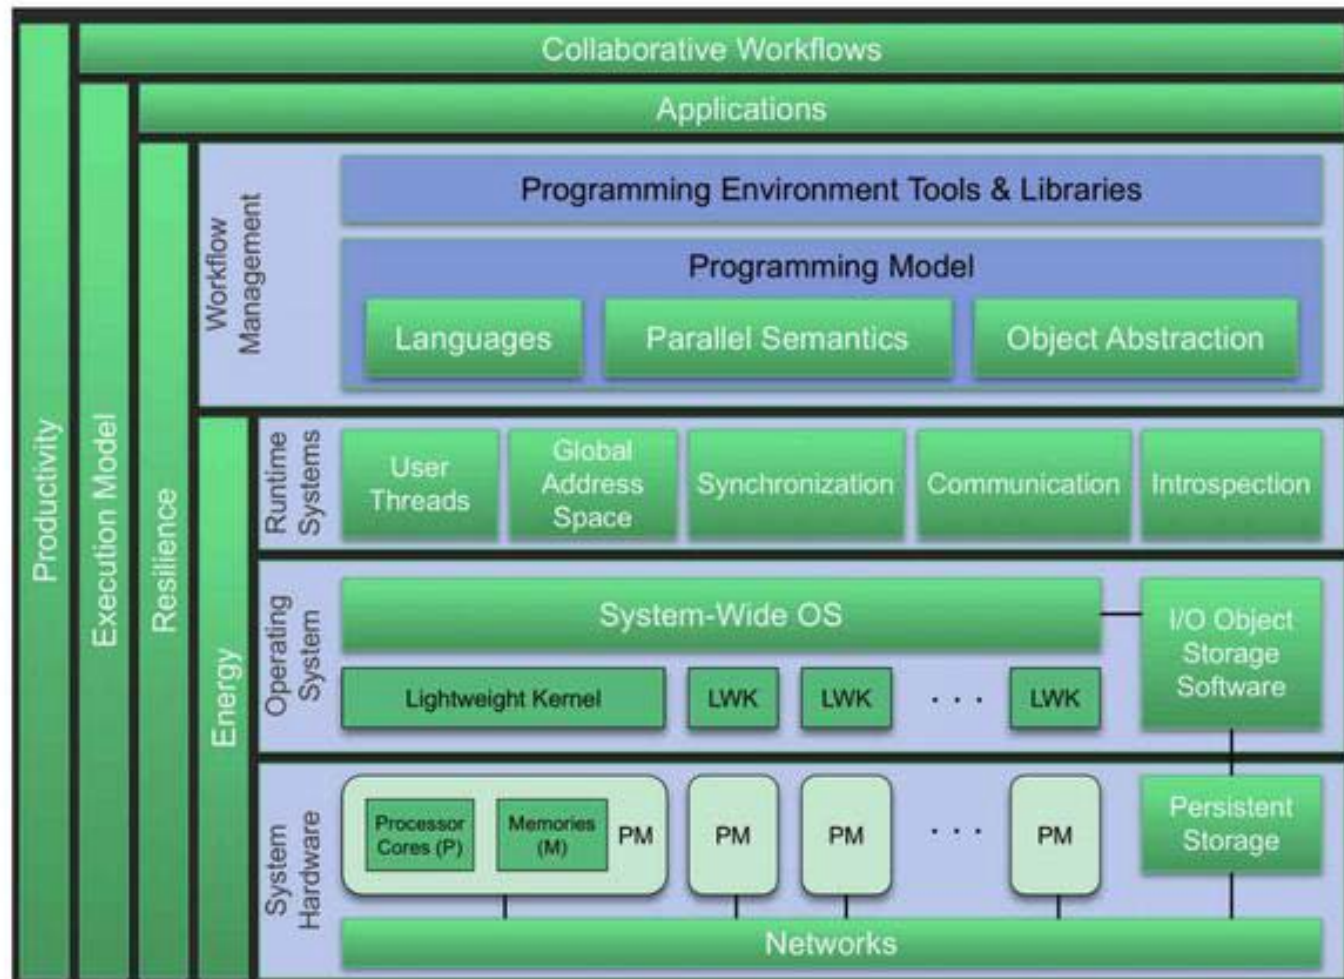
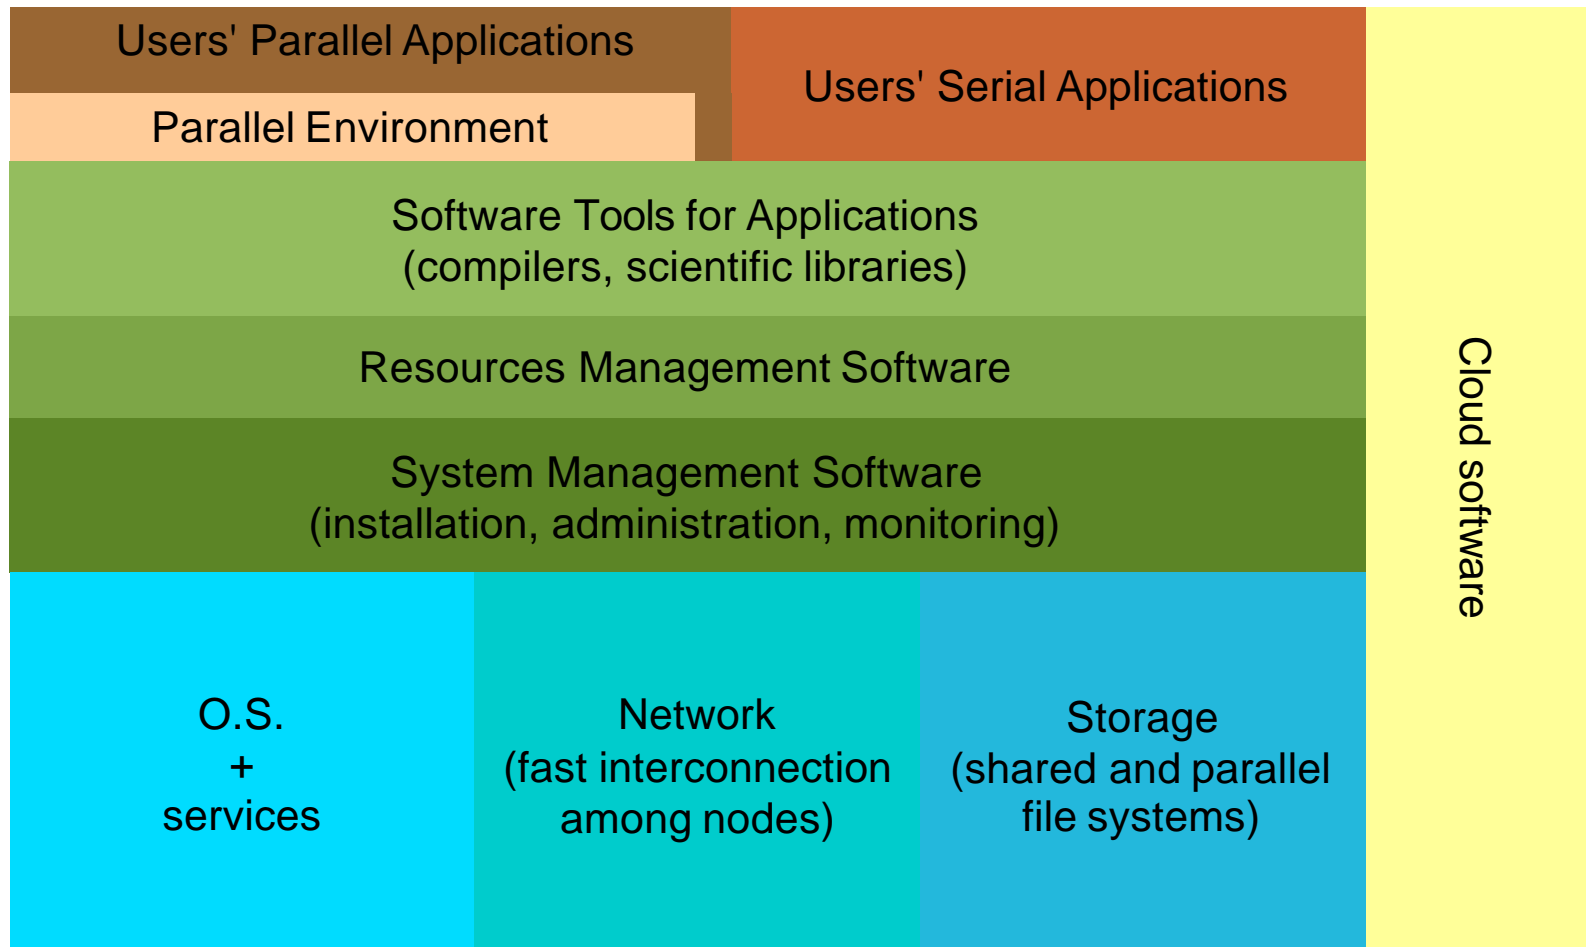
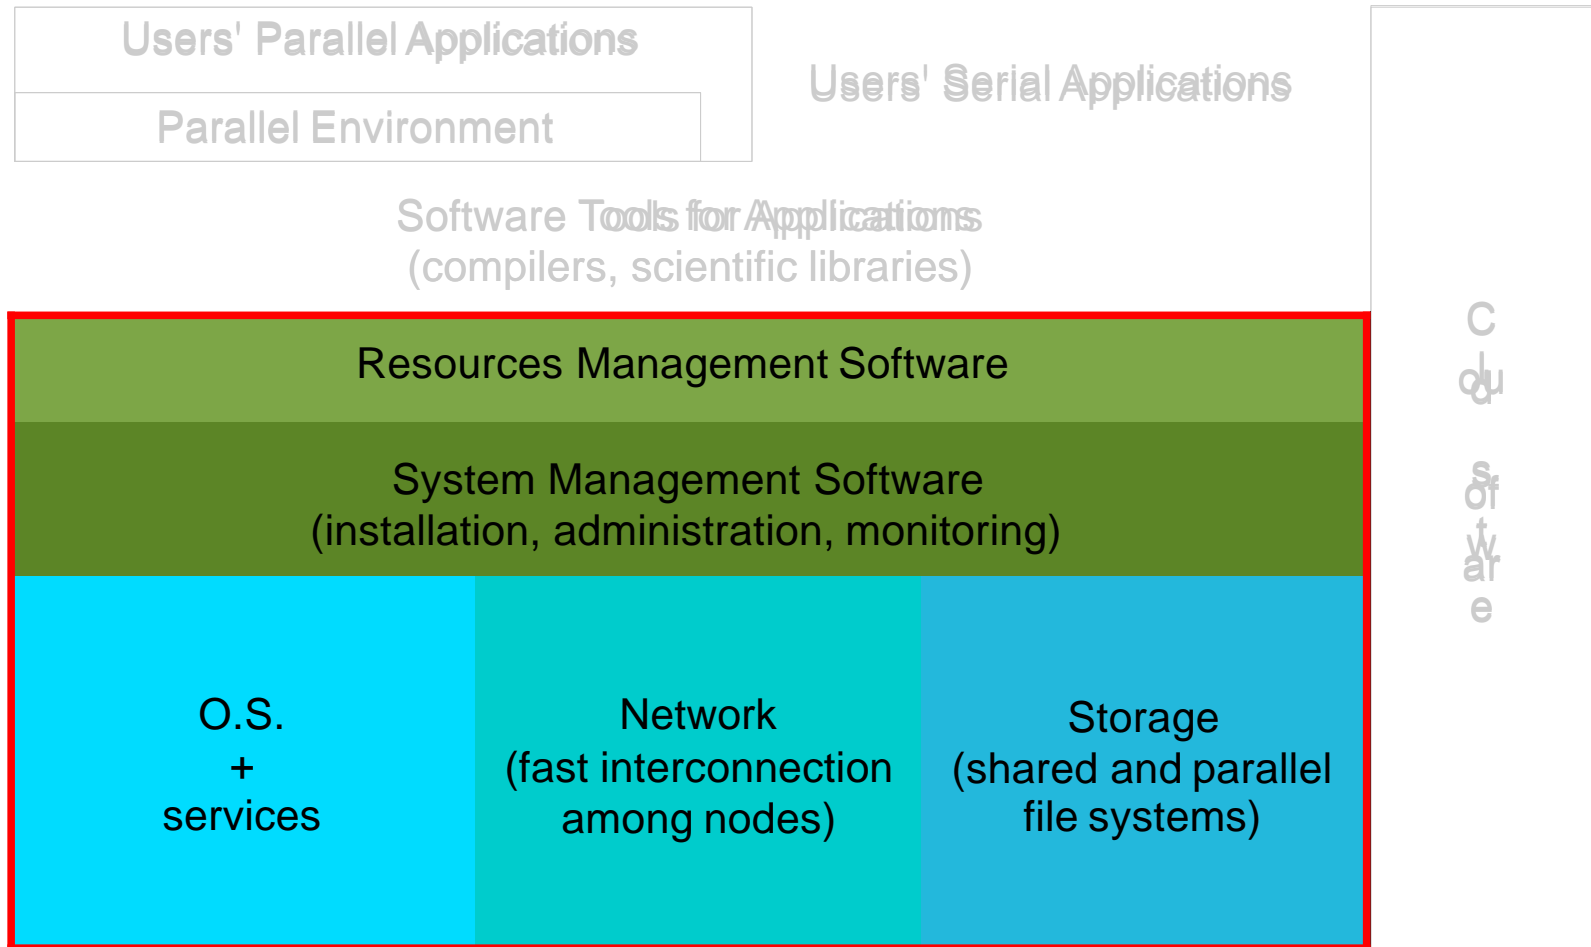


Figure 1.9 The system stack of a general supercomputer consists of a system hardware layer and several software layers. The first software layer is the operating system, encompassing both resource management and middleware to access input/output (I/O) channels. Higher software layers include runtime systems and workflow management.

A little bit simpler vision



The cluster middleware

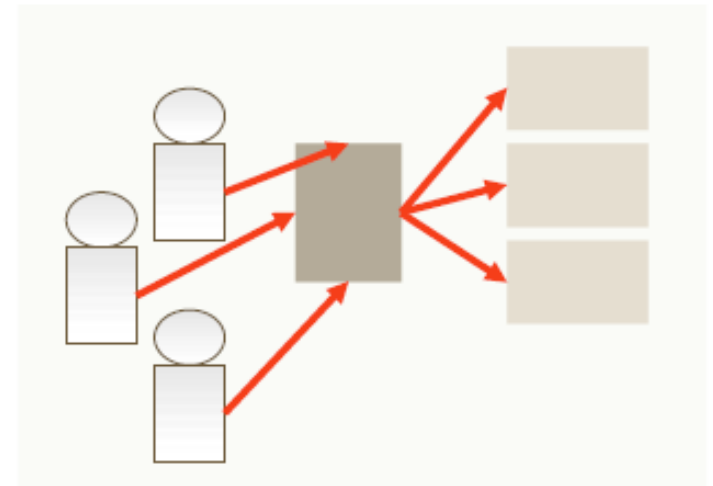
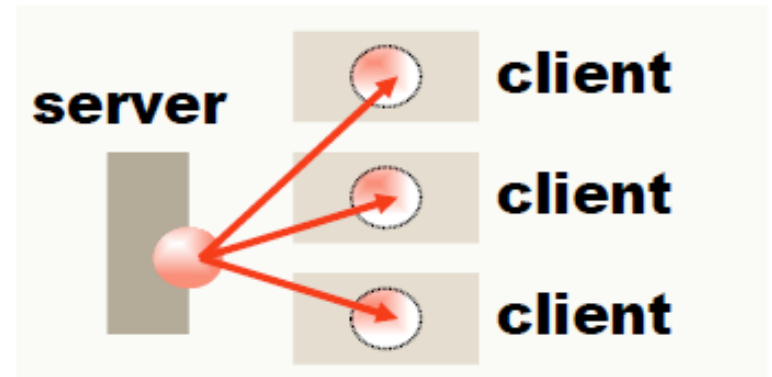


Cluster middleware design goals

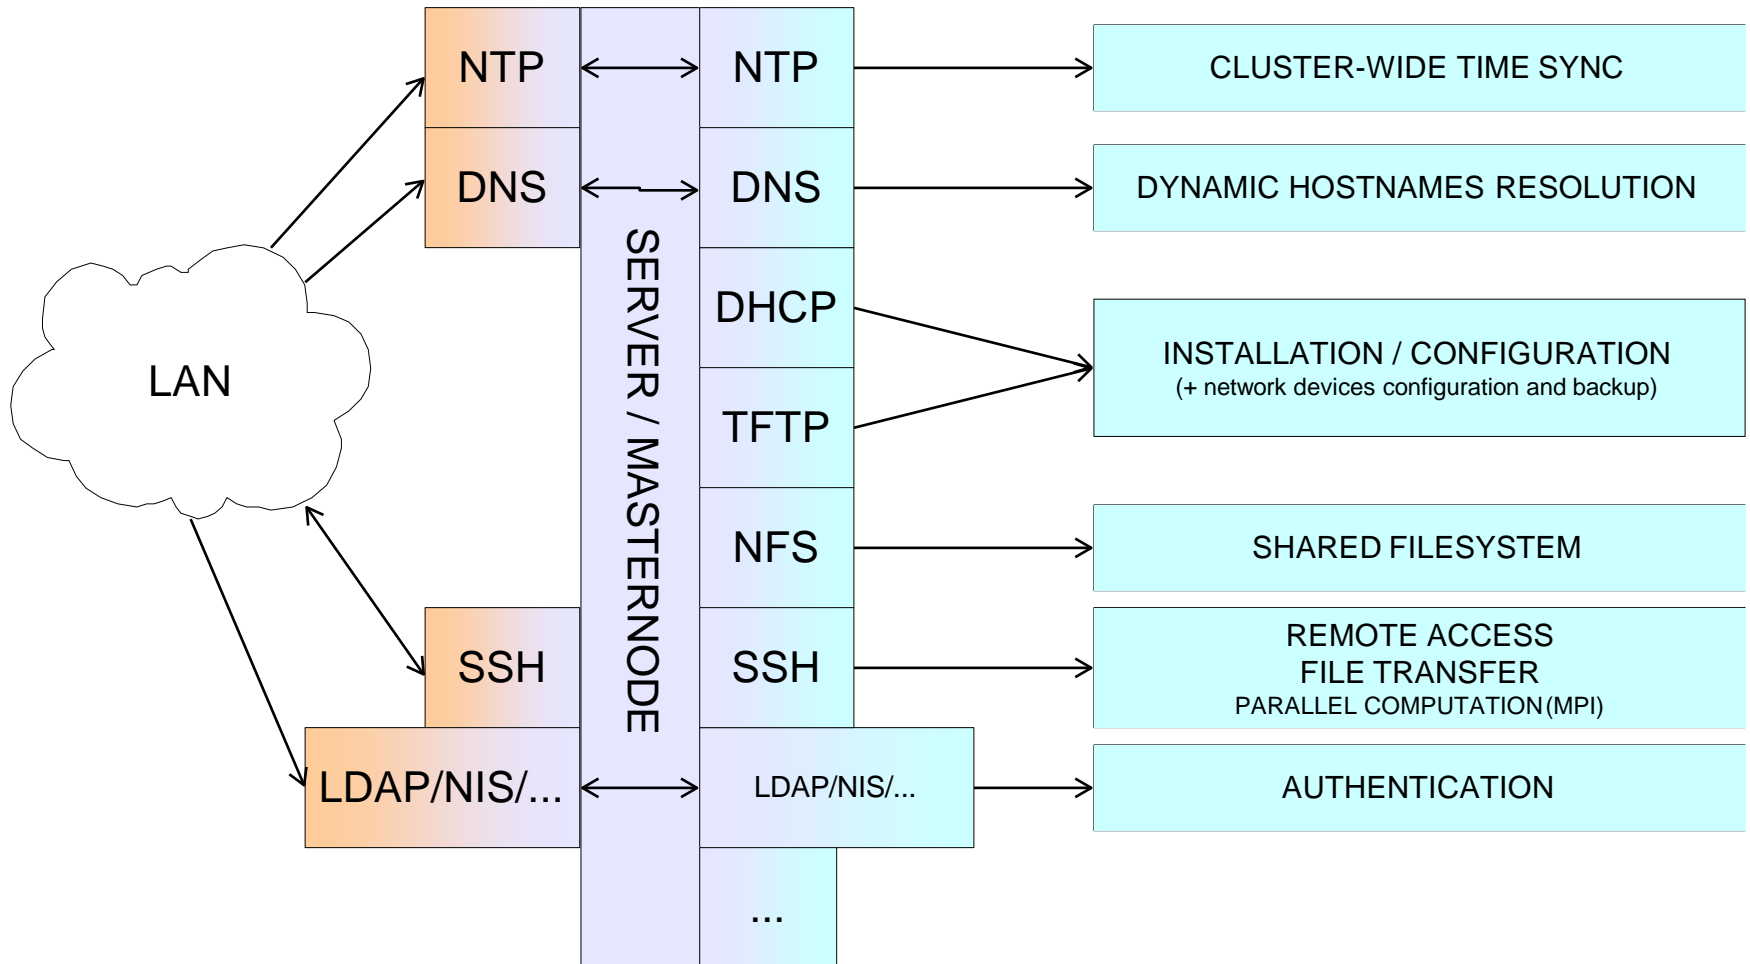
- Complete Transparency (Manageability)
 - Lets us see a single cluster system..
 - Single entry point: login ssh, software loading...
 - Unique storage for all nodes
- Scalable Performance:
 - Easy growth of cluster
- Enhanced Availability:
 - fault tolerant technologies
 - Automatic Recovery from failures

Cluster middleware

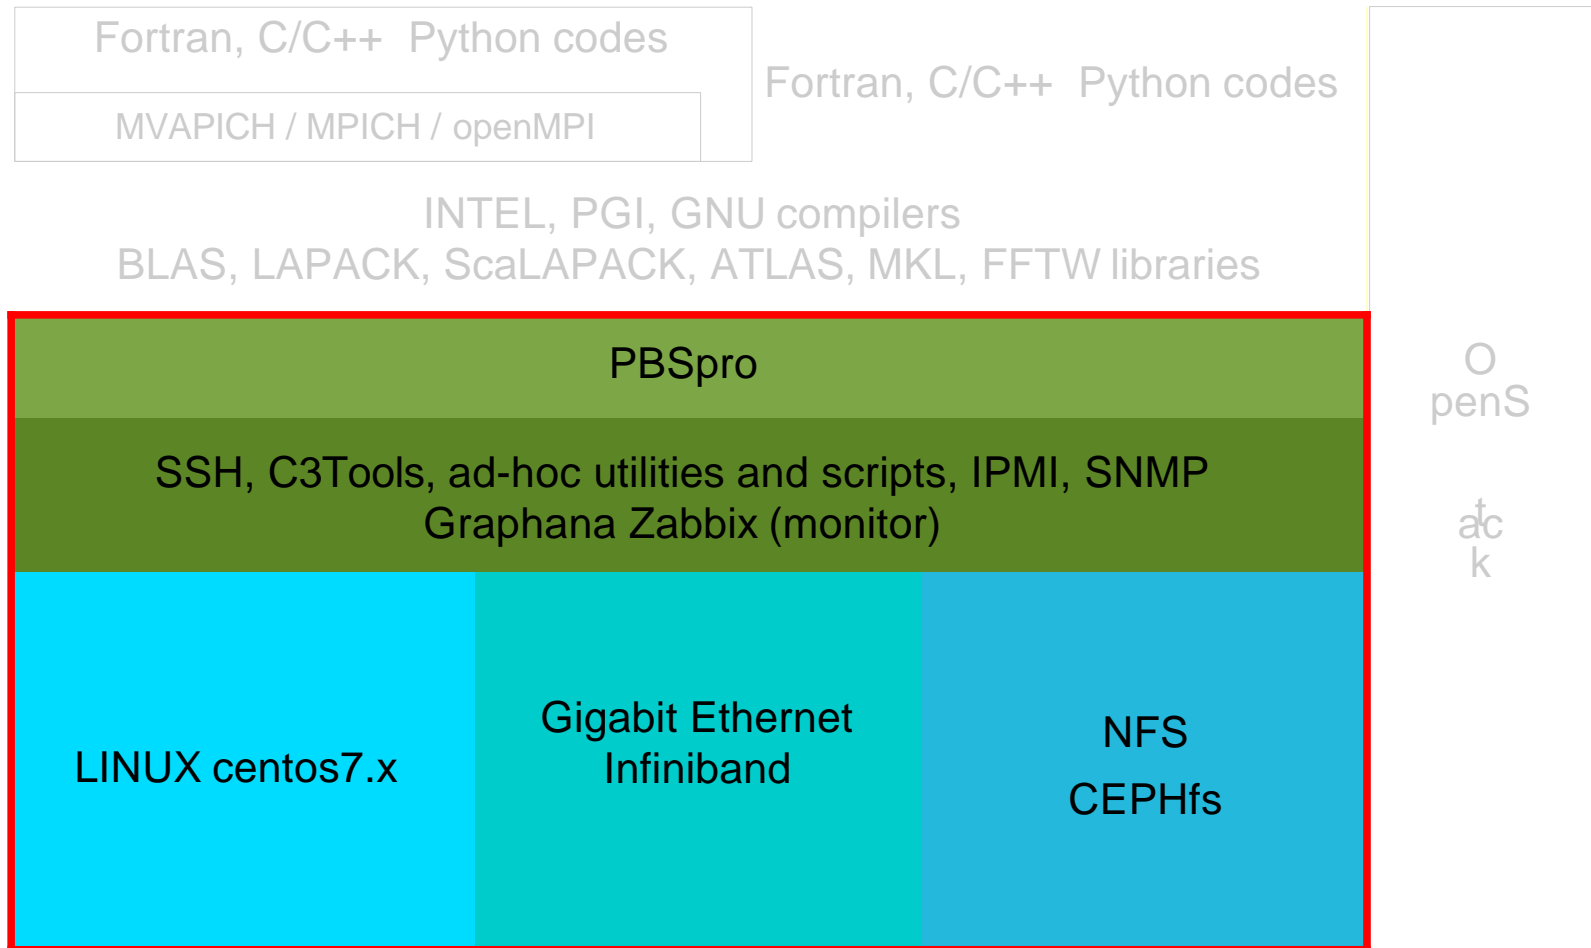
- Administration software:
 - user accounts
 - NTP/NFS/ etc...
- Resource management and scheduling software (LRMS)
 - Process distribution
 - Load balance
 - Job scheduling of multiple tasks



Cluster wide services



Middleware software used on ORFEO



Agenda

A first look of the software stack



Local resource manager: queue system

Scientific software

Compilers

Libraries

Resource Management Problem

- We have a pool of users and a pool of resources, then what?
 - some software that controls available resources
 - some other software that decides which application to execute based on available resources
 - some other software devoted to actually execute applications

What are we speaking about ?




REPLACE THE CAKE WITH HPC RESOURCE

Some definition

- **Batch Scheduler:** software responsible for scheduling the users' jobs on the cluster.
scheduling is the method by which work specified by some means is assigned to resources that complete the work
- **Resources Manager:** software that enable the jobs to connect the nodes and run.
- **Node** (aka Computing Node): computer used for its computational power.
- **Login/Master node:** it's through this node that the users will submit/launch/manage jobs.

Batch scheduler:




- Allocate resources for each applications with respect of their requirements and users' rights.

→ Satisfy users
response time, reliability

→ Satisfy admins
high resource utilization
efficiency, energy
management

Batch scheduler (2)




- Resource Management Layer
→ launching, cleaning, monitoring

Job Management Layer

- batch/interactive job
- backfilling
- scheduling
- suspend/Resume
- preemption
- dependencies
- resubmission
- advance reservation

Batch scheduler (3)



- Workload/Job Management
 - more complete job scheduling policies
 - Fairsharing, Quality of Service (QoS), SLA (Service Level Agreement), Energy Saving
 - Sometime a dedicated software

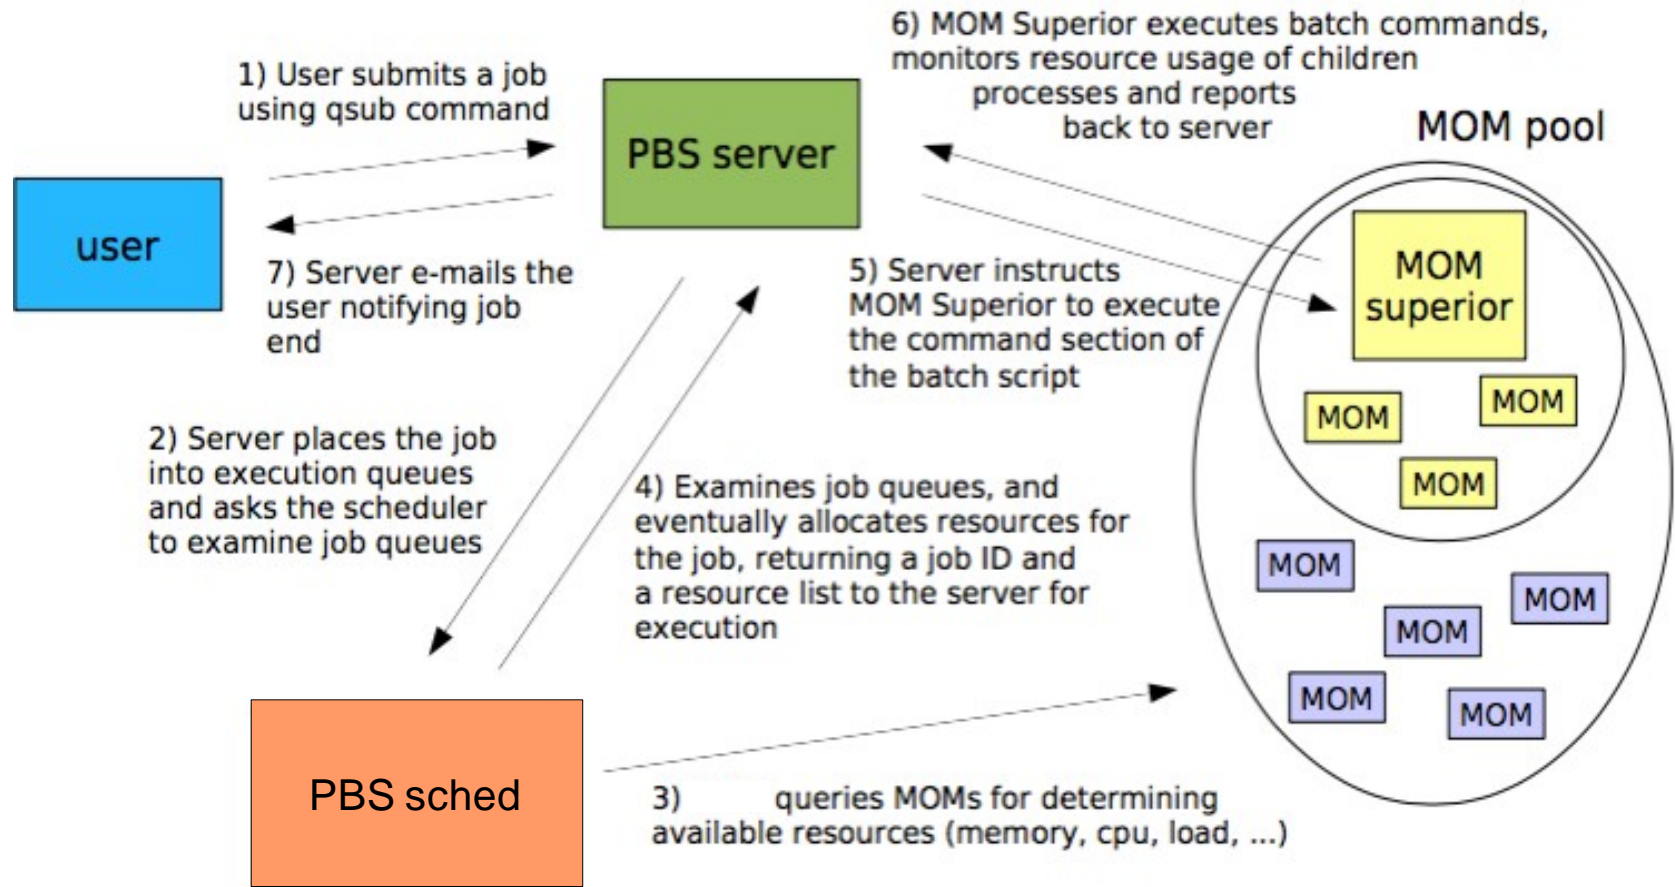
Main LRMS packages

- IBM LSF
 - commercial
- Univa Grid Engine (UGE)
 - Commercial originates from SGE
- PBSPRO
 - Portable Batch System Professional once commercial now open
 - Support is commercial
 - **Available on ORFEO**
- SLURM
 - Open source
 - Support is commercial

A job's life...

- The user describes the resources he needs in a shell script, the job file
- From the login node, the user submits the job to the queue system
- The system sends the job to the execution queue
- The job is executed on the compute node, without user intervention
- The results of the job are written in the folder specified by the user
- The queue system free the resources to get ready for the next execution

A PBSpro job's life



A pbs jobfile

<p>PBS directives, using the tag #PBS, describe the job requirements in terms of execution queue, number of nodes and cores, job name, walltime, etc.</p>	<pre>#!/bin/bash # This is an example script #PBS -q dssc #PBS -l nodes=1:ppn=2 #PBS -N myjob #PBS -l walltime=2:00:00</pre>
<p>The rest of the job is a standard shell script</p> <p>PBS "lands" user's home directory: it is important to change the directory to the one in which we want to run the job</p>	<pre>cd \$HOME/MyJobDir hostname pwd</pre>

Recap on LRMS

- LRMS is a fundamental tool in the HPC management:
 - User: know it well and you will almost run !
 - Sys. Adm.: know it well and you will keep your system busy..
- Many different choices
- Concepts are similar /commands sometime also (to help survive: <http://www.schedmd.com/slurmdocs/rosetta.pdf>, available on our repo)
- Key point is THE scheduler
 - Theoretically is almost all possible in resource scheduling with modern LRMS software to accommodate requests from users
 - Practically is almost impossible satisfy all your users (and/or communities)

Resource sharing policies is not at all a technical problem !

Agenda

A first look of the software stack



Local resource manager: queue system




Scientific software

Compilers

Libraries

HPC scientific Software layers (interleaved..)

- User's applications (both parallel and serial)
- Parallel Libraries&Tools
- Mathematical/Scientific Libraries
- I/O libraries
- Compilers



Who cares about scientific software ?

- End Users of HPC Software
 - Install and run HPC applications and tools
- HPC Application Teams
 - Manage third-party dependency libraries
- Package Developers
 - People who want to package their own software for distribution
- User support teams at HPC Centers
 - People who deploy software for users at large HPC sites

HPC software

- Not much standardization in HPC: every machine/app has a different software stack
 - This is done to get the best performance
- HPC frequently trades reuse and usability for performance
 - Reusing a piece of software frequently requires you to port it to many new platforms
- List of packages/combination can diverge...

Dependency Nightmare..

Scientific software: where is ?

- Generally available cluster-wide
- installed in /opt/cluster/software (or similar) and mounted read-only on the nodes via nfs
- Generally managed by **modules package**
- Several versions managed by some agreement

Module package (1)

- Modules allow to dynamically modify user environment
- Useful tool to track different version of installed software

Module package (2)

- A few useful commands

`module avail` – lists all available modules

`module list` – lists all loaded modules

`module load` – adds a module to your environment

`module unload` – removes a module from your environment

Module and environment

- Module command change on the fly the most important ENVIRONMENT VARIABLE for you
- PATH
- LD_LIBRARY_PATH

ORFEO situation

```
[cozzini@login ~]$ module avail
```

```
----- /opt/area/shared/modules/mpi -----  
openmpi/4.0.3/gnu/4.8.5 (D)    openmpi/4.0.3/gnu/9.3.0  
  
----- /opt/area/shared/modules/applications -----  
python/3.7.7/gnu/4.8.5    python/3.8.2/gnu/4.8.5  
  
----- /opt/area/shared/modules/utilities -----  
hwloc/2.2.0    numactl/2.0.13  
  
----- /opt/area/shared/modules/compilers -----  
cuda/11.0.3    gnu/9.3.0    intel/20.1
```

Where:

D: Default Module

Agenda

A first look of the software stack



Local resource manager: queue system



Scientific software



Compilers

Libraries

What does mean compiling ?

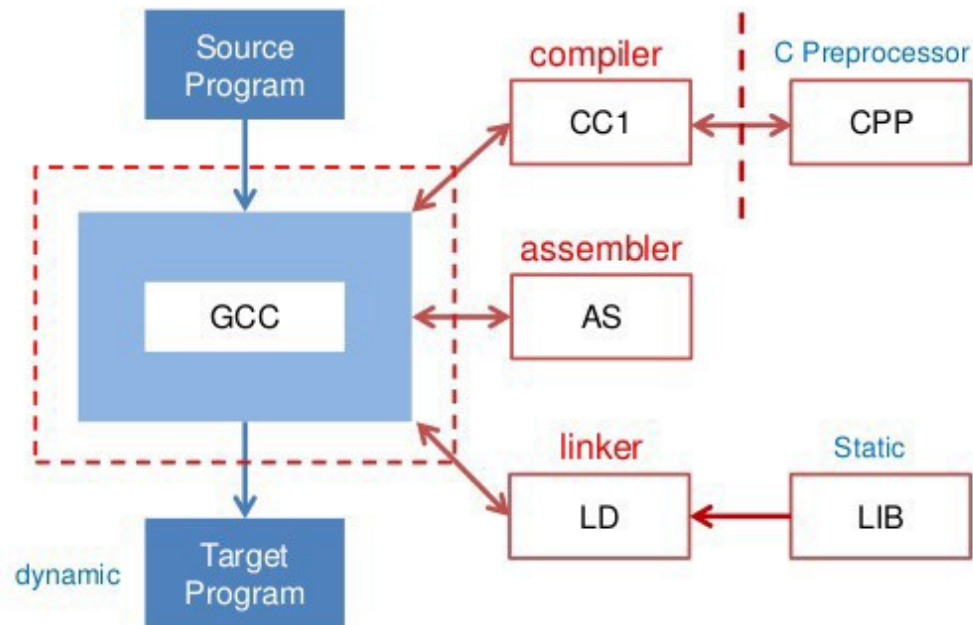
- A complex translation from high level language (C/Fortran...) to a stream of instructions..

Compiler

- Free : Gnu suite
 - Always available
 - Many different versions
 - Fundamental but some time lacks performance
- Commercial compilers
 - Intel suite :
 - A full software stack (includes libraries/ profiling /benchmarking tools MPI libraries
 - highly optimized
- PGI
 - Good compiler
 - Comes with some nice extension (openACC /Cuda Fortran)
 - Community edition available for free

Compiling with gcc..

GCC compiler



GCC is a collection that invokes compiler, assembler and linker...

What is available on ORFEO ?

```
[cozzini@login ~]$ module avail
```

```
----- /opt/area/shared/modules/mpi -----  
openmpi/4.0.3/gnu/4.8.5 (D)    openmpi/4.0.3/gnu/9.3.0
```

```
----- /opt/area/shared/modules/applications -----  
python/3.7.7/gnu/4.8.5    python/3.8.2/gnu/4.8.5
```

```
----- /opt/area/shared/modules/utilities -----  
hwloc/2.2.0    numactl/2.0.13
```

```
----- /opt/area/shared/modules/compilers -----  
cuda/11.0.3    gnu/9.3.0    intel/20.1
```

Where:

D: Default Module

Agenda

A first look of the software stack



Local resource manager: queue system



Scientific software



Compilers



Libraries

Scientific Libraries

- Plenty of them for many different tasks
- Dedicated lecture later during the course
- Today let us just focus on static vs dynamic libraries on basic system libraries

Static libraries: libfoo.o

- .a files are archives of .o files (object files)
- Linker includes needed parts of a static library in the output executable
- No need to find dependencies at runtime – only at build time.
- Can lead to large executables
- Often hard to build a completely static executable on modern systems.

Shared libraries: libfoo.so (Linux)

- More complex build semantics, typically handled by the build system
- Must be found by ld.so and loaded at runtime
- 2 main ways:
 - **LD_LIBRARY_PATH**: environment variable configured by user and/or module system
 - RPATH: paths embedded in executables and libraries, so that they know where to find their own dependencies.

All done !

A first look of the software stack



Local resource manager: queue system



Scientific software



Compilers



Libraries

