

# Optimization Avoid the avoidable

Luca Tornatore - I.N.A.F. 

“Foundation of HPC” course




**DATA SCIENCE &  
SCIENTIFIC COMPUTING**

2021-2022 @ Università di Trieste



Optimization


# Outline



Branches



Pipelines



Loops



Optimization

# Outline




Avoid the avoidable  
inefficiencies



Branches




Pipelines



Loops


For the purpose of setting-up an example,  
let's suppose that

- 1) we have a distribution of random data points on a 2D plane which we subdivide in sub-regions using a grid.



For the purpose of setting-up an example, let's suppose that


- 1) we have a distribution of random data points on a 2D plane which we subdivide in sub-regions using a grid.
- 2) for each point  $p$ , we want to select all the grid cells whose center is closer to  $p$  than a given radius  $r$ , and to perform some operations accordingly to our search result.






For the purpose of setting-up an example, let's suppose that

- 1) we have a distribution of random data points on a 2D plane which we subdivide in sub-regions using a grid.
- 2) for each point  $p$ , we want to select all the grid cells whose center is closer to  $p$  than a given radius  $r$ , and to perform some operations accordingly to our search result.



We may consider to  
use a nested loop  
like this one →

Is there anything  
you would change?



```
for(p = 0; p < Np; p++)  
  
    for(i = 0; i < Ng; i++)  
        for(j = 0; j < Ng; j++)  
            for(k = 0; k < Ng; k++)  
            {  
                dist = sqrt(  
                    pow(x[p] - (double)i/Ng - half_size, 2) +  
                    pow(y[p] - (double)j/Ng - half_size, 2) +  
                    pow(z[p] - (double)k/Ng - half_size, 2));  
  
                if(dist < R)  
                    do something;  
            }  
}
```



# | (1) Avoid expensive function calls



Some function calls are particularly expensive. Those include, among others, `sqrt()`, ...

Try to avoid them *if possible*.

```
for(p = 0; p < Np; p++)  
  
    for(i = 0; i < Ng; i++)  
        for(j = 0; j < Ng; j++)  
            for(k = 0; k < Ng; k++)  
                {  
                    dist2 = pow(x[p] - (double)i/Ng - half_size, 2) +  
                        pow(y[p] - (double)j/Ng - half_size, 2) +  
                        pow(z[p] - (double)k/Ng - half_size, 2));  
  
                    if(dist2 < R2)  
                        do something;  
                }
```





# | (1) Avoid expensive function calls



Some function calls are particularly expensive. Those include, among others, `sqrt()`, `pow()`, ...

Try to avoid them *if possible*.

```
for(p = 0; p < Np; p++)  
  
    for(i = 0; i < Ng; i++)  
        for(j = 0; j < Ng; j++)  
            for(k = 0; k < Ng; k++)  
                {  
                    dx = x[p] - (double)i/Ng - half_size;  
                    dy = y[p] - (double)j/Ng - half_size;  
                    dz = z[p] - (double)k/Ng - half_size;  
  
                    dist2 = dx*dx + dy*dy + dz*dz;  
                    if(dist2 < R2)  
                        do something;  
                }
```



# | (1) Avoid expensive function calls



Some function calls are particularly expensive. Those include, among others, `sqrt()`, `pow()`, floating point division, .. Try to avoid them *if possible*.

```
for(p = 0; p < Np; p++)  
  
    for(i = 0; i < Ng; i++)  
        for(j = 0; j < Ng; j++)  
            for(k = 0; k < Ng; k++)  
                {  
                    dx = x[p] - (double)i * Ng_inv - half_size;  
                    dy = y[p] - (double)j * Ng_inv - half_size;  
                    dz = z[p] - (double)k * Ng_inv - half_size;  
  
                    dist2 = dx*dx + dy*dy + dz*dz;  
                    if(dist2 < R2)  
                        do something with sqrt(dist2);  
                }
```



# | (1) Avoid expensive function calls



`(double)<i,j,k> * Ng_inv + half_size`

was performed  $N^3 + N^2 + N$  times, always returning the same values.

Hoisting would save

$N(N^2 + N + 1)$  **mul**, **add** and **mem** accesses.

You can do better pre-computing the relevant values:

```
double ijk[Ng];
for(i = 0; i < Ng; i++)
    ijk[i] = i * Ng_inv + half_size
```

```
for(p = 0; p < Np; p++)
```

```
    for(i = 0; i < Ng; i++)
        for(j = 0; j < Ng; j++)
            for(k = 0; k < Ng; k++)
            {
                dx = x[p] - (double)i * Ng_inv - half_size;
                dy = y[p] - (double)j * Ng_inv - half_size;
                dz = z[p] - (double)k * Ng_inv - half_size;

                dist2 = dx*dx + dy*dy + dz*dz;
                if(dist2 < R2)
                    do something with sqrt(dist2);
            }
```



## | (2) Hoisting of expressions

```
for(i = 0; i < Ng; i++) {  
    dx2 = x[p] - (double)i * Ng_inv - half_size;  
    dx2 = dx2*dx2;
```

```
    for(j = 0; j < Ng; j++) {  
        dy2 = y[p] - (double)j * Ng_inv - half_size;  
        dy2 = dy2*dy2;  
        dist2_xy = dx2 + dy2;
```

```
        for(k = 0; k < Ng; k++) {  
            dz = z[p] - (double)k * Ng_inv - half_size;  
            dist2 = dist2_xy + dz*dz;  
            if(dist2 < Rmax2)  
                do something with sqrt(dist2); } } }
```

$(\text{double})\langle i, j, k \rangle * Ng\_inv + half\_size$

was performed  $N^3 + N^2 + N$  times, always returning the same values.

Hoisting would save

$N(N^2 + N + 1)$  **mul**, **add** and **mem** accesses.



## | (2) *Hoisting* of expressions



You could do even better by pre-computing the relevant values:

```
double ijk[Ng];  
for(i = 0; i < Ng; i++)  
    ijk[i] = i * Ng_inv + half_size
```

```
for(i = 0; i < Ng; i++) {  
    dx2 = x[p] - Ng_inv[i] - half_size;  
    dx2 = dx2*dx2;  
  
    for(j = 0; j < Ng; j++) {  
        dy2 = y[p] - Ng_inv[j] - half_size;  
        dist2_xy = dx2 + dy2*dy2;  
  
        for(k = 0; k < Ng; k++) {  
            dz = z[p] - Ng_inv[k] - half_size;  
            dist2 = dist2_xy + dz*dz;  
            if(dist2 < Rmax2)  
                do something with sqrt(dist2); } } }
```



# | (3) Clarify the variables' scope



All these variables are very local, there's no need for them to have a wider scope.

That will help you in writing the code, and *may* help the compiler in optimizing the stack and perhaps the registers usage.

```
for(int i = 0; i < Ng; i++) {  
    double dx2 = x[p] - (double)i * Ng_inv - half_size;  
    dx2 *= dx2;  
  
    for(j = 0; j < Ng; j++) {  
        double dy2 = y[p] - (double)j * Ng_inv - half_size;  
        double dist2_xy = dx2 + dy2*dy2;  
  
        for(k = 0; k < Ng; k++) {  
            double dz = z[p] - (double)k * Ng_inv - half_size;  
            double dist2 = dist2_xy + dz*dz;  
  
            if(dist2 < Rmax2)  
                do something with sqrt(dist2); } } }
```





# | (4) Suggest what is important



These variables are often calculated and reused subsequently.

Keeping a register dedicated to them may be useful.

*Note: this is a suggestion, the compiler, after analyzing the code, may decide differently*

```
double register Ng_inv = 1.0 / Ng;
for(int i = 0; i < Ng; i++) {
    double dx2 = x[p] - (double)i * Ng_inv - half_size;
    dx2 *= dx2;

    for(j = 0; j < Ng; j++) {
        double dy2 = y[p] - (double)j * Ng_inv - half_size;
        dy2 *= dy2;
        double register dist2_xy = dx2 + dy2;

        for(k = 0; k < Ng; k++) {
            double register dz = z[p] - (double)k * Ng_inv - ...;
            double register dist2 = dist2_xy + dz*dz;

            if(dist2 < Rmax2)
                do something with sqrt(dist2); } } }
```



## | (5) Don't repeat unnecessary checks



Do you expect any great performance from this code ?

If not, why ?

```
char * find_char_in_string( char *string, char c )
{
    int i = pos;
    while ( i < strlen(string) )
        if( string[i] == c )
            break;
        else
            i++;

    if( i < strlen(string) )
        return &string[i];
    else
        return NULL;
}
```



## | (5) Don't repeat unnecessary checks



There are several details that dump the performance, i.e. the CPE, of this loop.

The one I want to draw your attention to is the repeated call to the `strlen()` function.

Do you expect the string to change while you are scanning it? I guess no, but the compiler does not know it and has no way to understand that by code analysis. Moreover, the memory pointed by `string` could be modified somewhere else between two iterations.

```
char * find_char_in_string( char *string, char c )
{
    int i = pos;
    while ( i < strlen(string) )
        if( string[i] == c )
            break;
        else
            i++;

    if( i < strlen(string) )
        return &string[i];
    else
        return NULL;
}
```



## | (5) Don't repeat unnecessary checks



This very simple change  
will save you a lot of CPE

```
char * find_char_in_string( char *string, char c )
{
    int i = pos;
    int len = strlen(string);
    while ( i < len )
        if( string[i] == c )
            break;
        else
            i++;

    if( i < strlen(string) )
        return &string[i];
    else
        return NULL;
}
```



## | (5) Don't repeat unnecessary checks



For a number of reasons, this version is even more efficient than the previous one.

Can you tell why ?

```
char * find_char_in_string( char *string, char c )
{
    char *pos = string;
    while( ( *pos != '\0' ) && ( *pos != c ) )
        pos++;

    return ( *pos == '\0' ? NULL : pos );
}
```



## (6) Avoid unnecessary memory references



This simple loop for a reduction of an array accumulates the partial results de-referencing the pointer `sum` at each iteration.

```
void reduce_vector( int n, double *array, double *sum )
{
    for ( int i = 0; i < n; i++ )
        *sum += array[i];
    return;
}
```





## (6) Avoid unnecessary memory references

This simple loop for a reduction of an array accumulates the partial results de-referencing the pointer `sum` multiple times.

( asm obtained with -O1)

```
void reduce_vector( int n, double *array, double
*sum )
{
    for ( int i = 0; i < n; i++ )
        *sum += array[i];
    return;
}
```

```
.L3:
    movsd xmm0, QWORD PTR [rdx]
    addsd xmm0, QWORD PTR [rax]
    movsd QWORD PTR [rdx], xmm0
    add    rax, 8
    cmp    rax, rcx
    jne     .L3
```


**movsd** *xmm0, value of \*sum*

**addsd** *xmm0, value of \*array*

**movsd** *address of sum, xmm0*

**add** *rax, 8* ( *array++* )

**cmp** *rax, n*



## (6) Avoid unnecessary memory references

Introducing a separated,  
local accumulator will  
save memory accesses

( asm obtained with -O1)

```
void reduce_vector( int n, double *array, double
*sum )
{
    double cum = 0;
    for ( int i = 0; i < n; i++ ) cum += array[i];
    *sum = cum;
    return;
}
```

```
.L11:
addsd xmm0, QWORD PTR [rax]
add    rax, 8
cmp    rax, rdx
jne    .L11
```

**addsd** xmm0, *value of \*array*

**add** rax, 8 (*array++*)

**cmp** rax, rdx (*array with end-of-array*)